Lumiere: A Space-Time Diffusion Model for Video Generation

Authors: Bar-Tal et al.

SIGGRAPH 2024, 259 citations

PRESENTED BY:
ANANTAPADMANAABHA PRASANNAKUMAR

Outline

- Motivation
- Method
- Evaluations
- Applications
- Societal Impact
- Limitations
- Conclusion

Motivation

- Restricted capability of existing models
  - Sensitive to error
  - Suffers from memory and computing constraints
  - Obtaining large-scale data is cumbersome
  - Training large-scale T2V is challenging

Motivation

- Employing temporal cascades is hindersome
  - Generates aggressively sub-sampled set of keyframes
  - TSR modules are constrained to fixed, small temporal context
  - Cascaded training suffers from domain gap

Method - Lumiere

- Utilizes Diffusion Probabilistic Models
  - Through denoising steps, trained to approximate data distribution
  - Starting from noise, the model obtains target distribution
- Incorporates additional guiding signals

Method - Pipeline

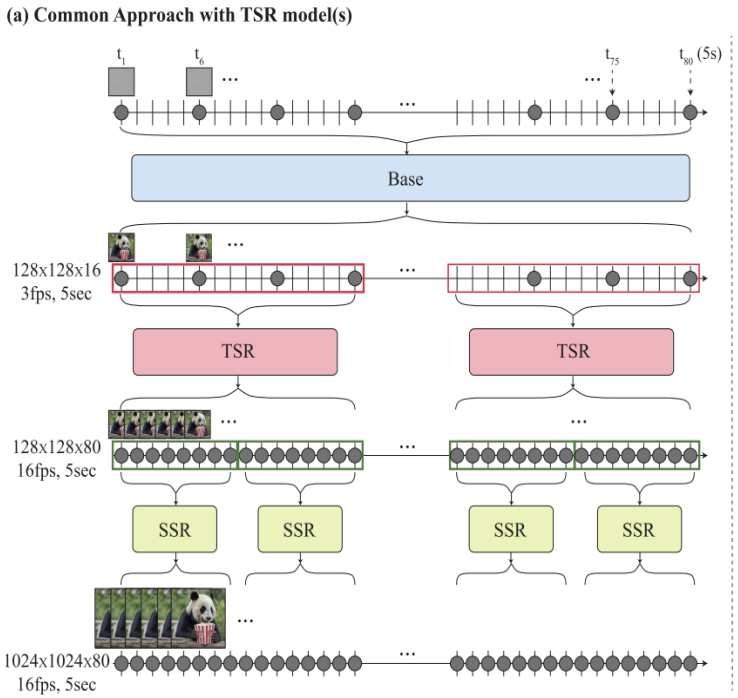

Common T2V Framework:

- Base Model
- Temporal Super-resolution Model (TSR)
- Spatial Super-resolution Model (SSR)

Method - Pipeline

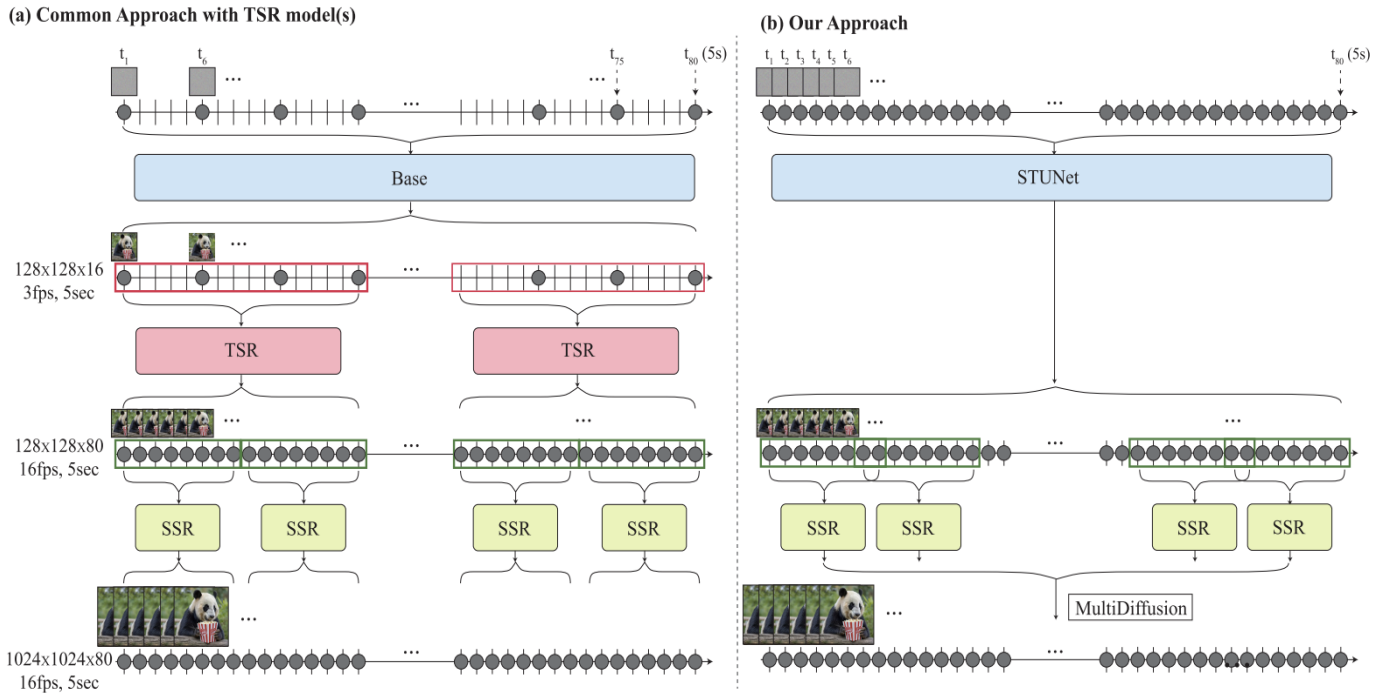

Lumiere Framework:

- Base Model
- Spatial Super-resolution Model (SSR)
- Multidiffusion

Lumiere: A Space-Time Diffusion Model for Video Generation, SIGGRAPH 2024, https://arxiv.org/pdf/2401.12945

Method - STUnet

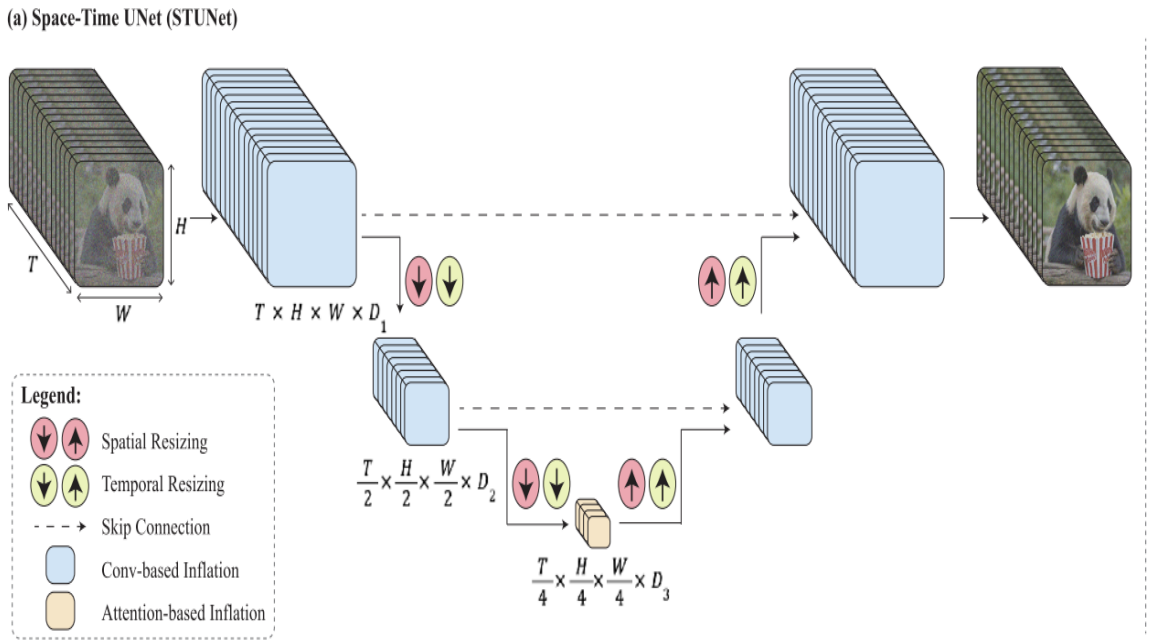

Employs traditional U-Net Model

Method - STUnet

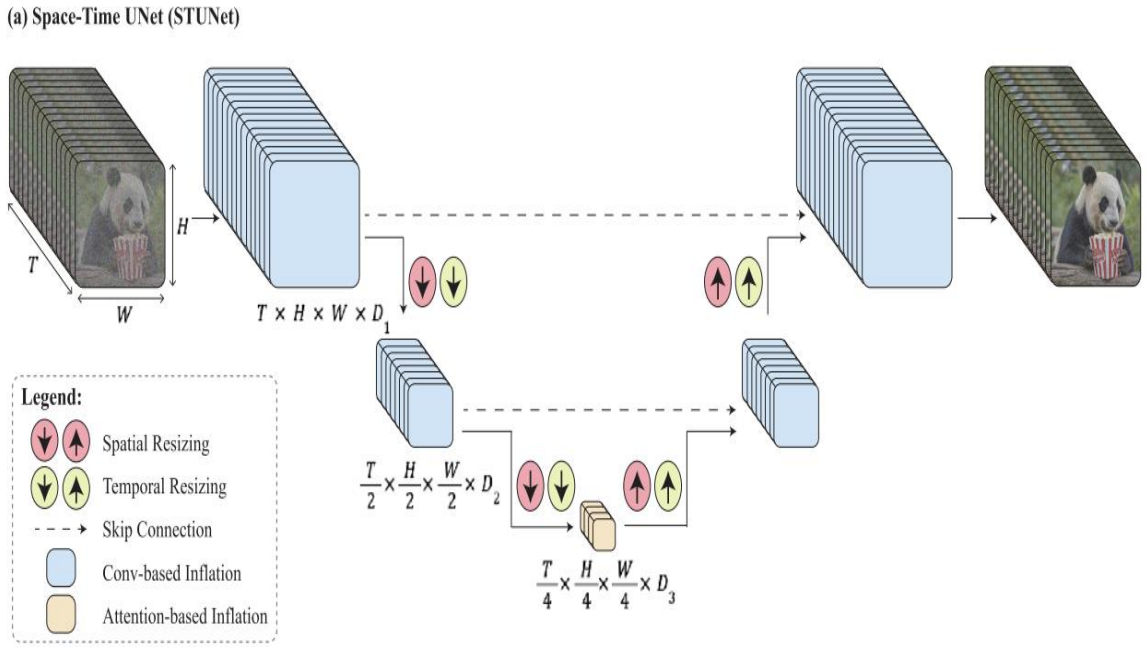

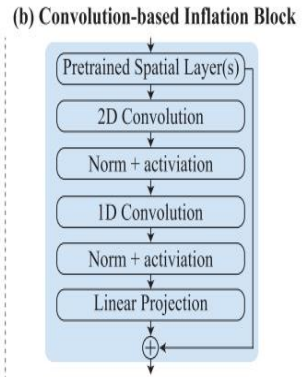

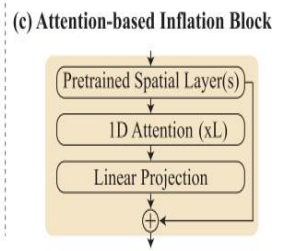

- Interleave temporal blocks to T2I Architecture
- Train only new parameters
- Low computational Overhead

Method - Multidiffusion

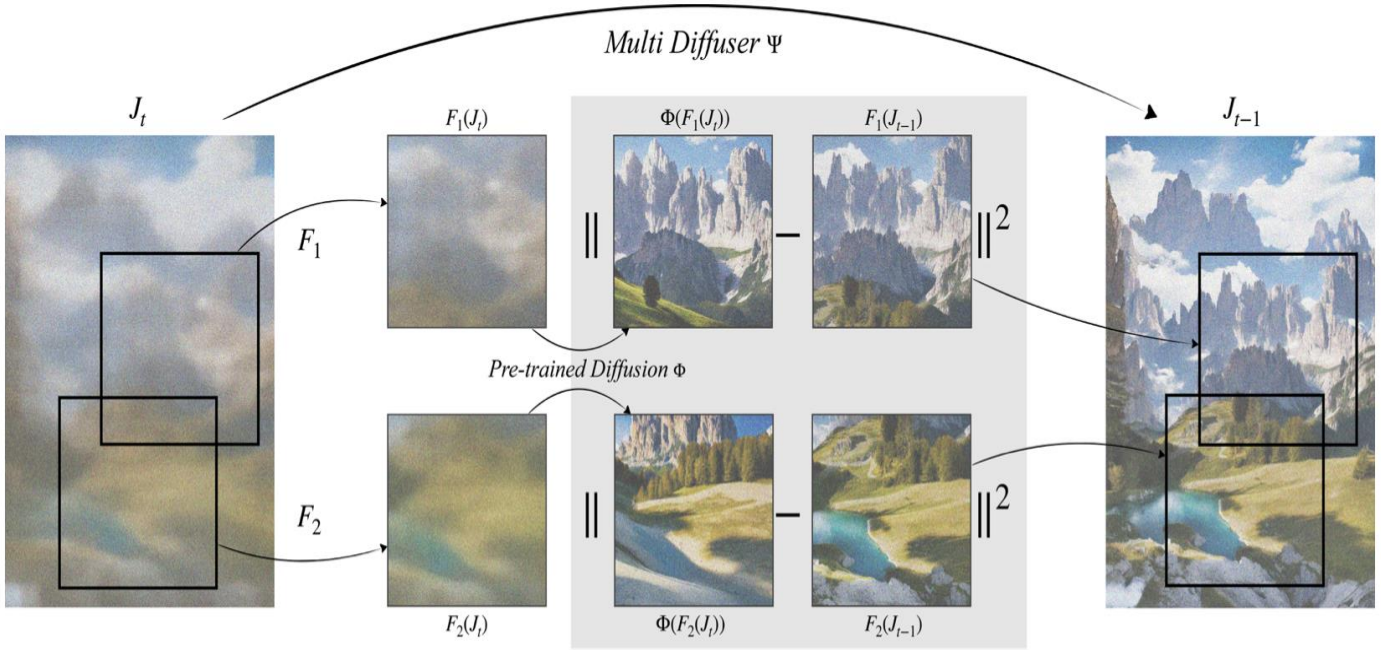

Multidiffusion Framework:

- New generation process from a pre-trained model
- Fuses inconsistent directions into one global denoising step

MultiDiffusion: Fusing Diffusion Paths for Controlled Image Generation, ICML 2023, https://arxiv.org/pdf/2302.08113

SSR with Multidiffusion

- An inflated SSR network can only operate on short videos
- Employ multidiffusion for smooth temporal transition
- Multidiffusion prevents temporal artifacts
  - Resolved by linearly combining video segments

SSR with Multidiffusion

Mathematically:

- At each generation step:
  - split noisy input video $J \in \mathbb{R}^{H \times W \times T \times 3}$ into $i \ldots N$ overlapping segments
  - Where $J_i \in \mathbb{R}^{H \times W \times T' \times 3}$ is the i^{th} segment
  - Temporal duration: $T' < T$
- To reconcile per-segment SSR predictions during denoising step:

$$\arg\min_{J'} \sum_{i=1}^{n} \|J' - \Phi(J_i)\|^2 .$$

Evaluation Setup

- Train T2V model on 30M videos with text prompts
  - Videos are 80 frames long at 16 fps
  - 109 text prompts
  - Base model dimension: 128 x 128 frames
  - SSR dimension: 1024 x 1024 frames

Zero-shot on UCF-101

Method	FVD ↓	IS ↑
MagicVideo (Zhou et al., 2022)	655.00	-
Emu Video (Girdhar et al., 2023)	606.20	42.70
Video LDM (Blattmann et al., 2023b)	550.61	33.45
Show-1 (Zhang et al., 2023a)	394.46	35.42
Make-A-Video (Singer et al., 2022)	367.23	33.00
PYoCo (Ge et al., 2023)	355.19	47.76
SVD (Blattmann et al., 2023a)	242.02	-
Lumiere (Ours)	332.49	37.54

User Study

- Two-alternative Forced Choice protocol Adopted
  - Randomly ordered pairs of videos are provided
  - 400 user judgments obtained
  - 109 prompts were utilized
  - Fixed random seed
  - Spatial and Temporal alignment

User Study

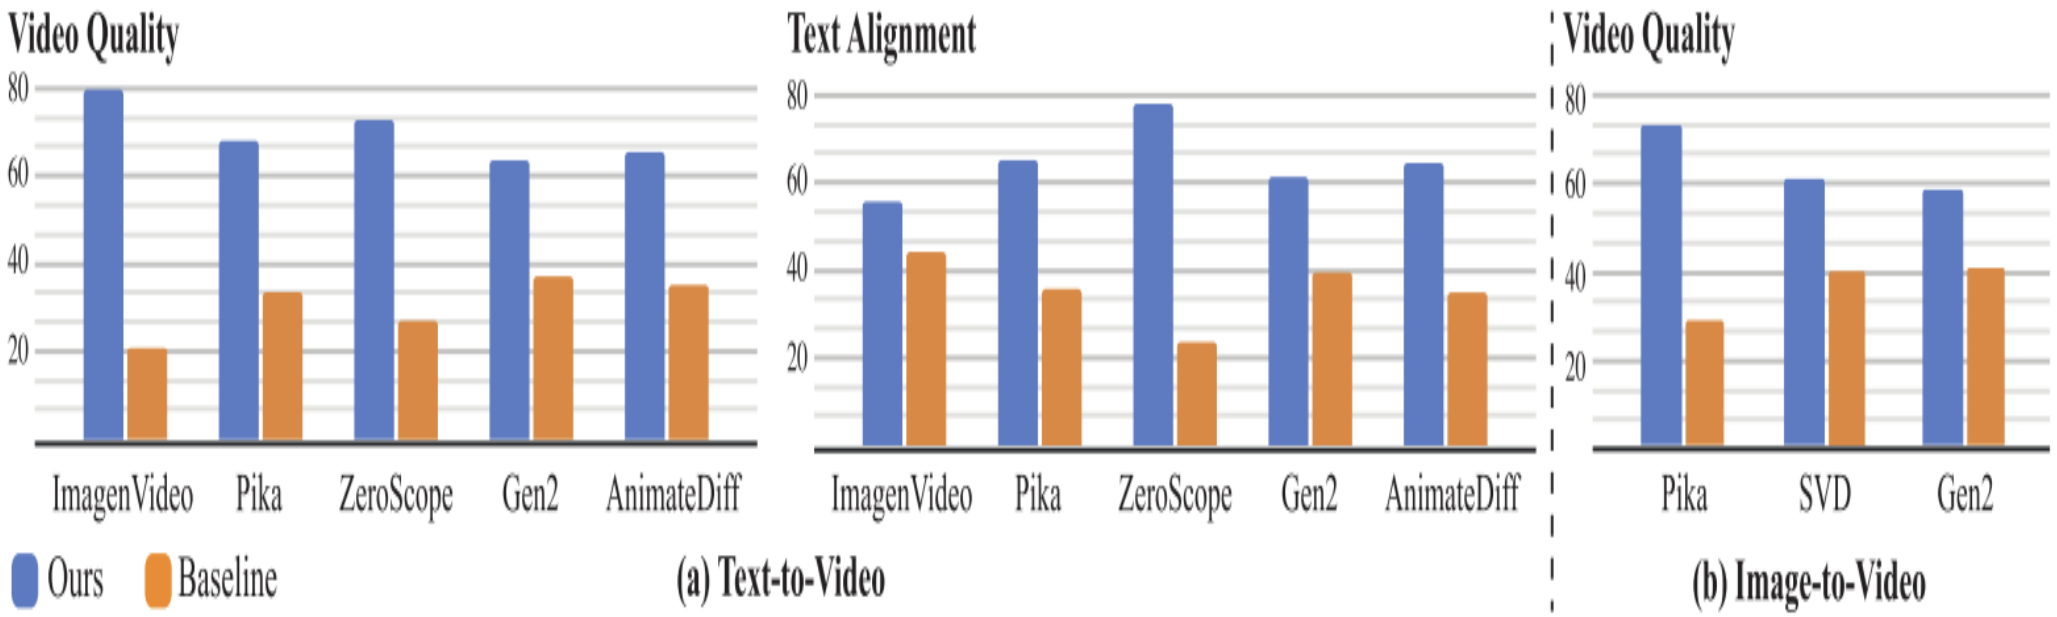
Video Quality
80
60
40
20
ImagenVideo
Pika
ZeroScope
Gen2
AnimateDiff
Ours
Baseline
Text Alignment
80
60
40
20
ImagenVideo
Pika
ZeroScope
Gen2
AnimateDiff
(a) Text-to-Video
Video Quality
80
60
40
20
Pika
SVD
Gen2
(b) Image-to-Video

User Study

Left video

Right video

Qualitative Evaluation

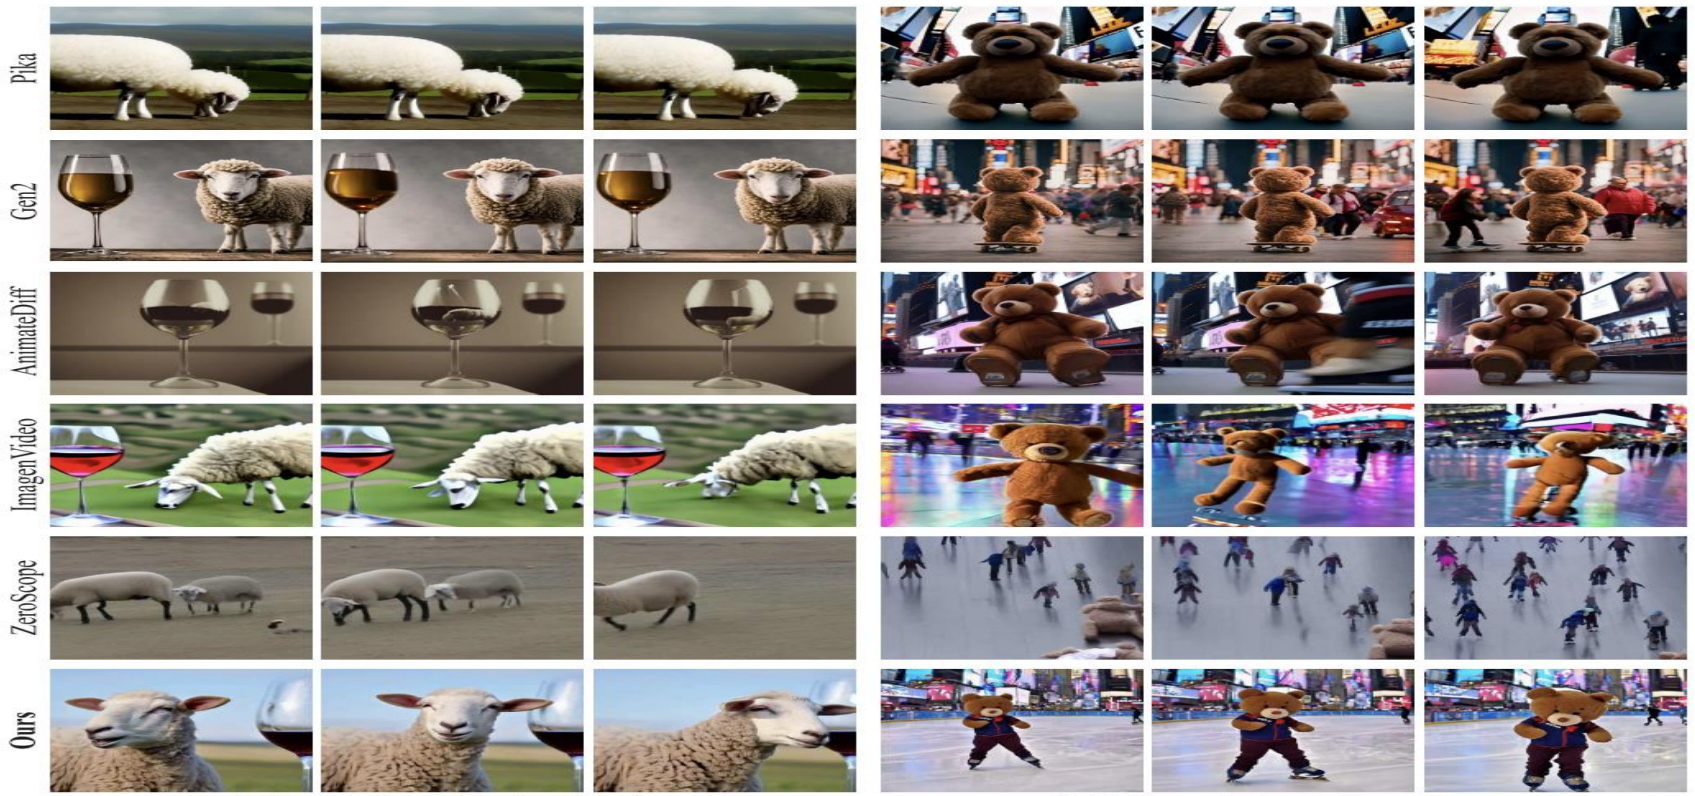

Temporal Consistency

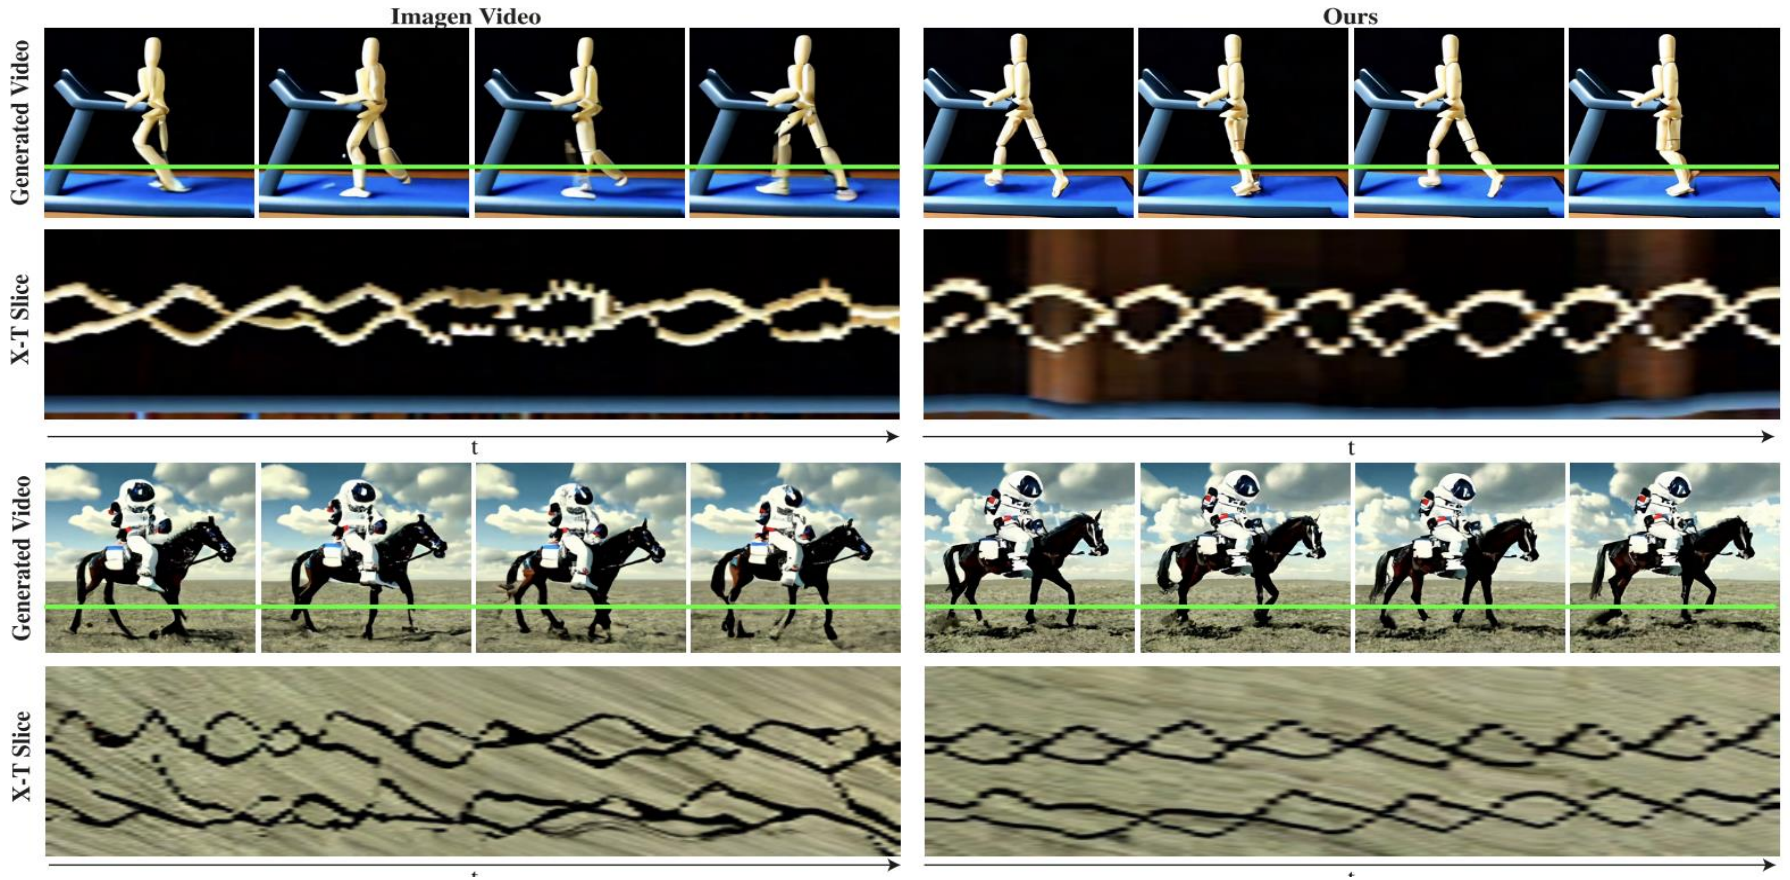

Ablation - Initialization

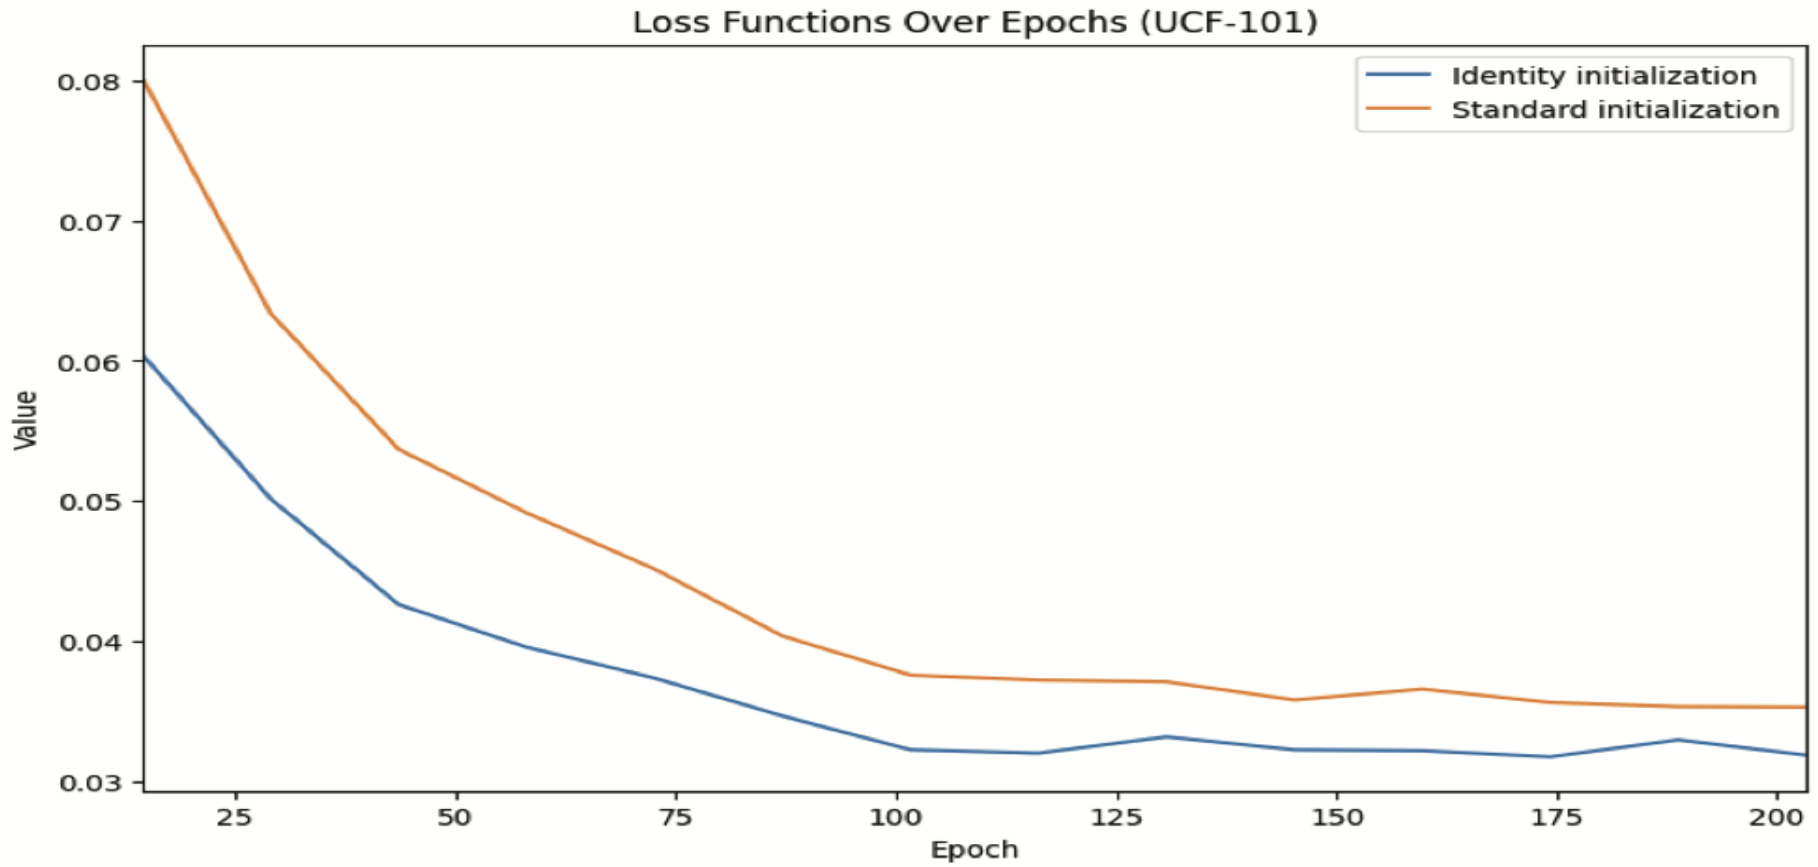

Ablation - Initialization Schemes

Ablation - Multidiffusion

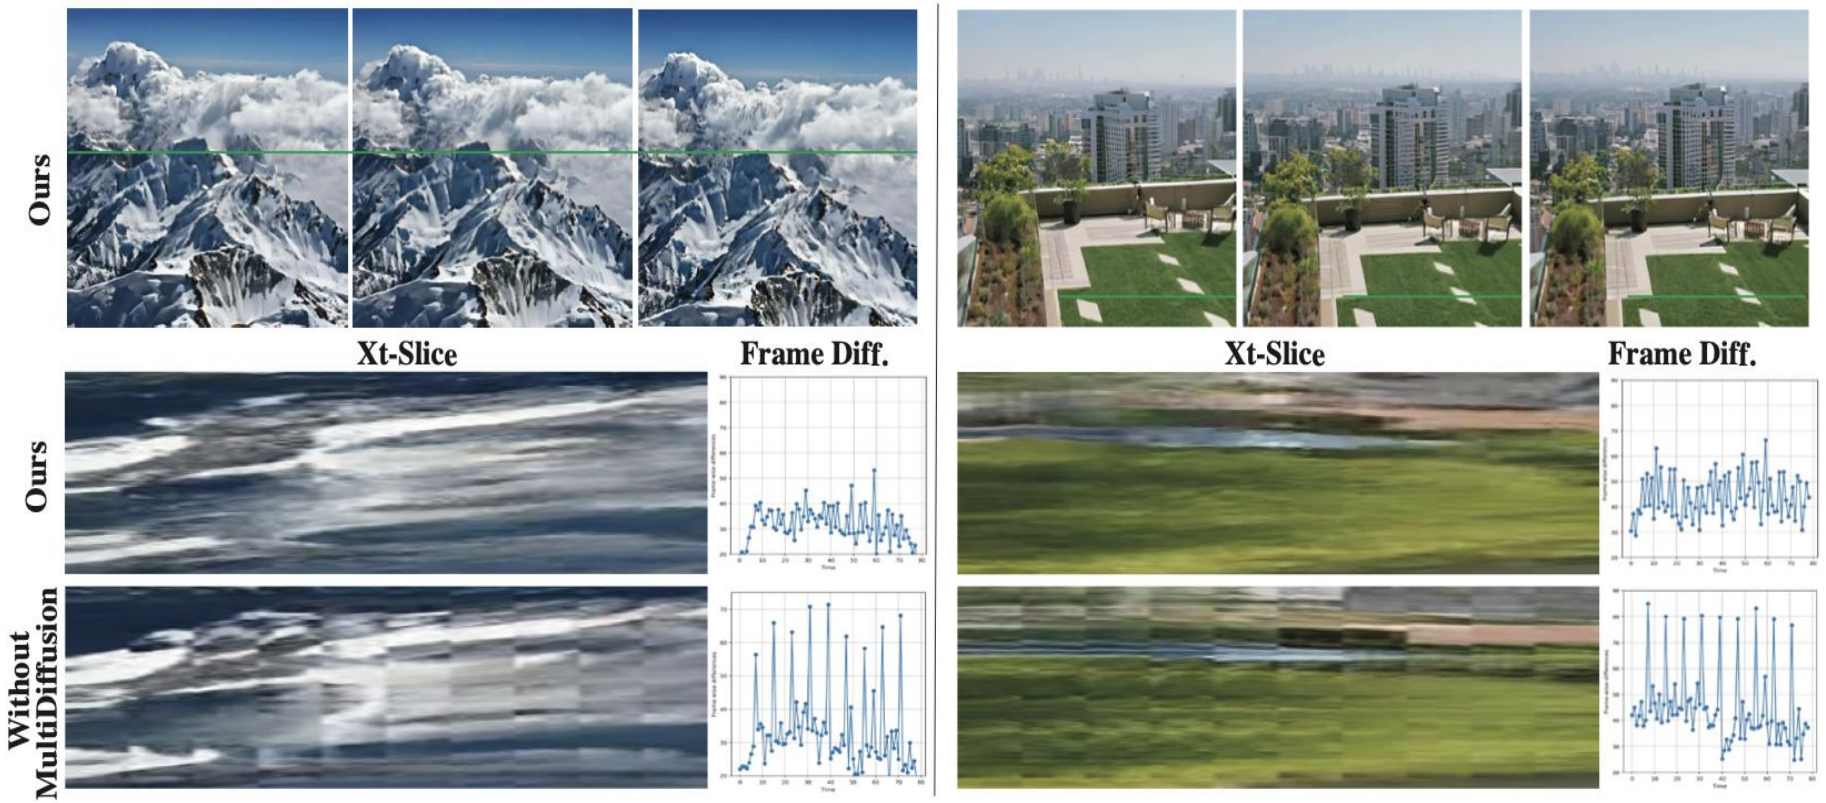

Applications – Video Editing

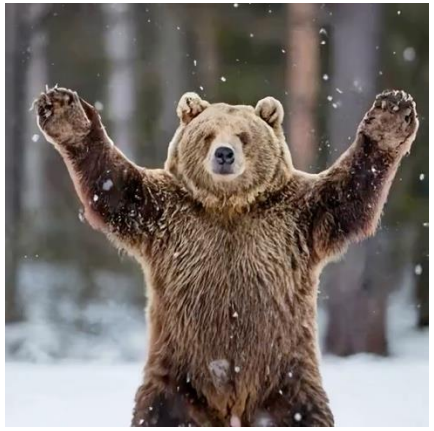

SDEDIT: GUIDED IMAGE SYNTHESIS AND EDITING WITH STOCHASTIC DIFFERENTIAL EQUATIONS, Jan 2022, https://arxiv.org/pdf/2108.01073

Application – Stylized Generation

- Pre-trained T2I weights remain fixed
- Newly added temporal layers are trained
- Linear interpolation between fixed and fine-tuned T2I weights
  - $W_{interpolate} = \alpha \cdot \ W_{style} + (1 - \alpha) \cdot W_{orig}$
  - Where $\alpha \in [0.5, 1]$

Application – Stylized Generation

Vector art styles

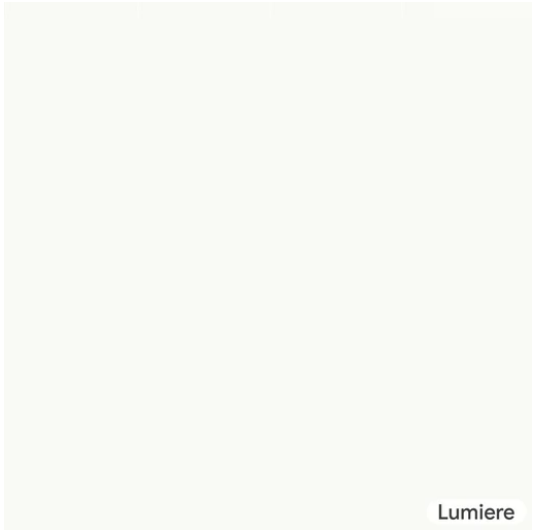

Application – Stylized Generation

Realistic styles

Application – Conditional Generation

- Model conditioned on additional input signals
  - Noisy video $J \in \mathbb{R}^{H \times W \times T \times 3}$
  - Text prompt
  - Masked conditioning video $C \in \mathbb{R}^{H \times W \times T \times 3}$
  - Masked video $M \in \mathbb{R}^{H \times W \times T \times 1}$

- Concatenated Tensor $\langle J, C, M \rangle = \mathbb{R}^{T \times H \times W \times 7}$

Application – I2V

Application - Inpainting

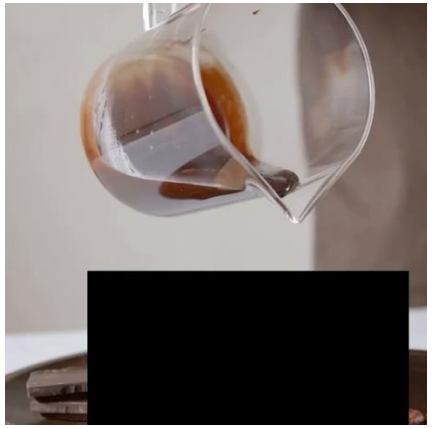

Application - Cinemagraphs

Societal Impact

- Risk of misuse
  - Tools for detecting biases and malicious use cases
  - To ensure safe and fair use

Limitations

- The model cannot generate videos
  - Multiple shots
  - Transition between scenes
- The model operates in pixel space

Conclusion

- Presents a novel T2V framework
  - Built on a pre-trained T2I model
  - Introduces space-time U-Net Architecture
  - Utilizes Multidiffusion framework
- Demonstrates state-of-the-art generation results
- Showcases applicability to various downstream tasks

Thank you